



# Neural ODE solvers

Michał Stypułkowski  
University of Wrocław

Szkoła Letnia AI Tech, Gdańsk, 20-24 maja 2022



**Fundusze Europejskie**  
Polska Cyfrowa



**Rzeczpospolita  
Polska**



**Unia Europejska**  
Europejski Fundusz  
Rozwoju Regionalnego



Projekt współfinansowany ze środków Unii Europejskiej w ramach Europejskiego Funduszu Rozwoju Regionalnego  
Program Operacyjny Polska Cyfrowa na lata 2014-2020.

Oś priorytetowa nr 3 „Cyfrowe kompetencje społeczeństwa”, działanie nr 3.2 „Innowacyjne rozwiązania na rzecz aktywizacji cyfrowej”.

Tytuł projektu: „Akademia Innowacyjnych Zastosowań Technologii Cyfrowych (AI Tech)”.



# Neural ODE solvers

# Introduction



# Differential equations

“Since Newton, mankind has come to realize that the **laws of physics** are always expressed in the language of **differential equations**.”

- Steven Strogatz

# Differential equations - overview

1. Ordinary differential equations (ODEs):
  - Equation with a function of **one independent variable**.
  - Often describing **change of function in time**.
  - Some of ODEs have closed-form solutions.

# Differential equations - overview

1. Ordinary differential equations (ODEs):
  - Equation with a function of **one independent variable**.
  - Often describing **change of function in time**.
  - Some of ODEs have closed-form solutions.
  
2. Partial differential equations (PDEs):
  - Equation with a function of **at least two independent variables**.
  - Usually no analytical solutions.

# Differential equations - examples

1. Ordinary differential equations (ODEs):
  - Pandemic model.
  - Predator-prey equations.
  - Newton's law of cooling/heating.
2. Partial differential equations (PDEs):
  - Diffusion process.
  - Fluid dynamics.
  - Black–Scholes equation.

# Differential equations - examples

## 1. Ordinary differential equations (ODEs):

- Pandemic model.
- Predator-prey equations.
- Newton's law of cooling/heating.



**FUN!**

## 2. Partial differential equations (PDEs):

- Diffusion process.
- Fluid dynamics.
- Black–Scholes equation.



# Differential equations - examples

## 1. Ordinary differential equations (ODEs):

- Pandemic model.
- Predator-prey equations.
- Newton's law of cooling/heating.



**FUN!**

## 2. Partial differential equations (PDEs):

- Diffusion process.
- Fluid dynamics.
- Black–Scholes equation.



**MORE FUN!**

But way more difficult :(

# Numerical ODE solvers

# Initial value problem

We face an **initial value problem** of the form:

$$x'(t) = f(t, x)$$

$$x(t_0) = x_0$$

We will often call function  $f$  a **dynamic function**.

# Taylor series

Using the Taylor series of a real infinitely differentiable function  $x$  at point  $t + h$  we can get:

$$x(t + h) = x(t) + hx'(t) + \frac{1}{2!}h^2x''(t) + \frac{1}{3!}h^3x'''(t) + \frac{1}{4!}h^4x^{(4)}(t) + \dots$$

# Taylor series

Using the Taylor series of a real infinitely differentiable function  $x$  at point  $t + h$  we can get:

$$x(t + h) = x(t) + hx'(t) + \frac{1}{2!}h^2x''(t) + \frac{1}{3!}h^3x'''(t) + \frac{1}{4!}h^4x^{(4)}(t) + \dots$$

In other words, when we know the function  $x$  and its value at time  $t$ , we can calculate its value **small step in the future** using **infinite** number of derivatives!

# Taylor series

Using the Taylor series of a real infinitely differentiable function  $x$  at point  $t + h$  we can get:

$$x(t + h) = x(t) + hx'(t) + \frac{1}{2!}h^2x''(t) + \frac{1}{3!}h^3x'''(t) + \frac{1}{4!}h^4x^{(4)}(t) + \dots$$

In other words, when we know the function  $x$  and its value at time  $t$ , we can calculate its value **small step in the future** using **infinite** number of derivatives!

**LET US APPROXIMATE!**



# Euler method

The simplest (1 order) method based on the Taylor series.

We can approximate the value of a  $x(t + h)$ :

$$x(t + h) \approx x(t) + hx'(t) = x(t) + hf(t, x)$$

# Euler method

The simplest (1 order) method based on the Taylor series.

We can approximate the value of a  $x(t + h)$ :

$$x(t + h) \approx x(t) + hx'(t) = x(t) + hf(t, x)$$

In practice, we use the following formula to update  $x_n$ :

$$x_{n+1} = x_n + hf(t, x_n)$$

# Runge-Kutta of order 4 (RK4)

This method is based on the Taylor series up to 4th derivative. The formula is:

$$x(t + h) \approx x(t) + \frac{1}{6}(F_1 + 2F_2 + 2F_3 + F_4)$$

where

$$F_1 = hf(t, x),$$

$$F_2 = hf\left(t + \frac{1}{2}h, x + \frac{1}{2}F_1\right),$$


$$F_3 = hf\left(t + \frac{1}{2}h, x + \frac{1}{2}F_2\right),$$

$$F_4 = hf(t + h, x + F_3).$$

# Unknown dynamic function

Let us now change the perspective.

Suppose we **do not know** the dynamic function but instead we are **given data points** from some trajectory.




# Unknown dynamic function

Let us now change the perspective.

Suppose we **do not know** the dynamic function but instead we are **given data points** from some trajectory.

We want to find a dynamic function that fits the data points best.




# Unknown dynamic function

Let us now change the perspective.

Suppose we **do not know** the dynamic function but instead we are **given data points** from some trajectory.

We want to find a dynamic function that fits the data points best.

**PARAMETRIZE IT!**





Neural network as an ODE


# Sequence of transformations

Models such as:

- Residual networks
- RNNs
- Normalizing flows

use a combination of transformations:

$$x_{t+1} = x_t + f(x_t, \theta_t)$$




# Sequence of transformations

Models such as:

- Residual networks
- RNNs
- Normalizing flows

use a combination of transformations:

$$x_{t+1} = x_t + f(x_t, \theta_t)$$



# Sequence of transformations


Models such as:

- Residual networks
- RNNs
- Normalizing flows

use a combination of transformations:

$$x_{t+1} = x_t + f(x_t, \theta_t)$$

This can be seen as an **Euler discretization!**



# More layers...



AI Memes for Artificially Intelligent Teens

@ai\_memes

...

Flight attendant: Is there a doctor on this flight?

Dad: \*nudging me\* that should've been you


Me: Not now Dad

Dad: Not asking for a DL practitioner to help, are they?

Me: Dad, there's a medical emergency happening right now

Dad: Go and see if "adding more layers" helps

Mor



ns

...

r on this flight?

e been you




oner to help, are they?

gency happening right


layers" helps



# Mor



- Unique optimum: global/local.




Schematic of a logistic regression classifier.

ns


You

The guy she tells you not to worry about



- Multiple local optima
- In high dimensions possibly


Deep neural network



Mor

Gentlemen  
overgene  
VC-Dimer  
is too high  
experts a  
structural  
Rework o  
make the  
unbiased  
soft marg


S  
M  
LA



$x_m$


Schematic of a logistic regression classifier.

The guy she tells you  
not to worry about



- Multiple local optima
- In high dimensions possibly

Deep neural network




out


Mor

Gentlemen  
overgeneralize  
VC-Dimension  
is too high  
experts are  
structural  
Rework or  
make the  
unbiased  
soft margin

S  
M  
L




Why she tells you  
worry about



local optima  
dimensions possibly  
Deep neural network

layer 1 hidden layer 2 hidden layer 3







Mor

Gentlemen  
overengineered  
VC-Dimer  
is too high  
experts a  
structural  
Rework o  
make the  
unbiased  
soft marg

S  
M  
LA



local optima  
dimensions possibly  
Deep neural network



# Continuous latent dynamics

If we add **infinite number of layers** and take **infinitely small steps**, we end up with a continuous dynamics on the hidden states:

$$\frac{dx(t)}{dt} = f(x(t), t, \theta)$$



# Continuous latent dynamics

If we add **infinite number of layers** and take **infinitely small steps**, we end up with a continuous dynamics on the hidden states:

$$\frac{dx(t)}{dt} = f(x(t), t, \theta)$$

Note, that the right side of the equation is a neural network.


# Continuous latent dynamics



## Example - MNIST classifier

Let us consider a problem of classifying MNIST digits.

We can define a dynamic function between input (image) and output (class label):




**4**




# Example - MNIST classifier

Let us consider a problem of classifying MNIST digits.

We can define a dynamic function between input (image) and output (class label):




$$x(0) = x_0$$

$$\frac{dx(t)}{dt} = f(x(t), t, \theta)$$


**4**

$$x(T) = y$$

## Example - MNIST classifier



$$x(0) = x_0$$

$$\frac{dx(t)}{dt} = f(x(t), t, \theta)$$



**4**

$$x(T) = y$$

Once we found the dynamic function, we can classify image using:

$$x(T) = x(0) + \int_0^T f(x(t), t, \theta) dt$$

# Dynamic function optimization

We can easily\* train our neural net using standard numerical methods (such as RK4).


Let us assume we want to optimize some loss function  $L$  :

$$L(\mathbf{z}(t_1)) = L \left( \mathbf{z}(t_0) + \int_{t_0}^{t_1} f(\mathbf{z}(t), t, \theta) dt \right)$$

# Dynamic function optimization

We can easily\* train our neural net using standard numerical methods (such as RK4).

Let us assume we want to optimize some loss function  $L$ :




# Dynamic function optimization

We can easily\* train our neural net using standard numerical methods (such as RK4).

Let us assume we want to optimize some loss function  $L$ :

**\*We need to remember every step!**



# Adjoint method

To make an optimization step, we need gradients  $\frac{dL}{d\theta}$ .

# Adjoint method

To make an optimization step, we need gradients  $\frac{dL}{d\theta}$ .

First, let us determine how the loss depends on the hidden state  $\mathbf{z}(t)$ .

Let us define **adjoint** as:

$$\mathbf{a}(t) = \frac{\partial L}{\partial \mathbf{z}(t)}$$

# Adjoint method

To make an optimization step, we need gradients  $\frac{dL}{d\theta}$ .

First, let us determine how the loss depends on the hidden state  $\mathbf{z}(t)$ .

Let us define **adjoint** as:

$$\mathbf{a}(t) = \frac{\partial L}{\partial \mathbf{z}(t)}$$

It turns out that the adjoint follows its own dynamic function:

$$\frac{d\mathbf{a}(t)}{dt} = -\mathbf{a}(t)^\top \frac{\partial f(\mathbf{z}(t), t, \theta)}{\partial \mathbf{z}}$$



# Adjoint method

To make an optimization step, we need gradients  $\frac{dL}{d\theta}$ .

First, let us determine how the loss depends on the hidden state  $\mathbf{z}(t)$ .

Let us define **adjoint** as:


$$\mathbf{a}(t) = \frac{\partial L}{\partial \mathbf{z}(t)}$$

It turns out that the adjoint follows its own dynamic function:

$$\frac{d\mathbf{a}(t)}{dt} = -\mathbf{a}(t)^\top \frac{\partial f(\mathbf{z}(t), t, \theta)}{\partial \mathbf{z}}$$

**One more pass through an ODE solver!**

# Adjoint method



# Adjoint method

Finally, we can calculate derivatives of the loss function with respect to model's parameters:

$$\frac{dL}{d\theta} = - \int_{t_1}^{t_0} \mathbf{a}(t)^\top \frac{\partial f(\mathbf{z}(t), t, \theta)}{\partial \theta} dt$$

# Adjoint method

Finally, we can calculate derivatives of the loss function with respect to model's parameters:

$$\frac{dL}{d\theta} = - \int_{t_1}^{t_0} \mathbf{a}(t)^\top \frac{\partial f(\mathbf{z}(t), t, \theta)}{\partial \theta} dt$$

Couple of remarks:

- It works with **O(1) memory!**

# Adjoint method

Finally, we can calculate derivatives of the loss function with respect to model's parameters:

$$\frac{dL}{d\theta} = - \int_{t_1}^{t_0} \mathbf{a}(t)^\top \frac{\partial f(\mathbf{z}(t), t, \theta)}{\partial \theta} dt$$

Couple of remarks:

- It works with **O(1) memory!**
- We can use it only on **imperfect optimization** problems (optimums are bad)!

# Adjoint method


Finally, we can calculate derivatives of the loss function with respect to model's parameters:

$$\frac{dL}{d\theta} = - \int_{t_1}^{t_0} \mathbf{a}(t)^\top \frac{\partial f(\mathbf{z}(t), t, \theta)}{\partial \theta} dt$$

Couple of remarks:

- It works with **O(1) memory!**
- We can use it only on **imperfect optimization** problems (optimums are bad)!
- For **less complicated models**, we should avoid this method due to the **slower training!**

# Adjoint method



Theory



torchdiffeq

# Continuous Normalizing Flow (CNF)




# Normalizing flows


## Inference

$$x \sim \hat{p}_X$$
$$z = f(x)$$

Data space  $\mathcal{X}$





Latent space  $\mathcal{Z}$



## Generation

$$z \sim p_Z$$
$$x = f^{-1}(z)$$



# Normalizing flows

We can express  $p_X$  (likelihood of the data) using **change of variable formula**:

$$p_X(x) = p_Z(f(x)) \left| \det \left( \frac{\partial f(x)}{\partial x^T} \right) \right|$$

We need  $f$  to be a bijection.

# Normalizing flows

We can express  $p_X$  (likelihood of the data) using **change of variable formula**:

$$p_X(x) = p_Z(f(x)) \left| \det \left( \frac{\partial f(x)}{\partial x^T} \right) \right|$$

We need  $f$  to be a bijection. But also:

- Easily invertible.
- Simple form of the determinant of the Jacobian.

# Real NVP

Let  $f = f_n \circ f_{n-1} \circ \dots \circ f_1$ , where  $f_i$  is defined as:

$$y_1 = x_1$$

$$y_2 = x_2 \odot \exp(s(x_1)) + t(x_1)$$

where  $(x_1, x_2)$  is some partition of the input.

# Real NVP

Let  $f = f_n \circ f_{n-1} \circ \dots \circ f_1$ , where  $f_i$  is defined as:

$$y_1 = x_1$$

$$y_2 = x_2 \odot \exp(s(x_1)) + t(x_1)$$

where  $(x_1, x_2)$  is some partition of the input.

- Easily invertible!
- Simple form of the determinant of the Jacobian:

$$\det \left( \frac{\partial f_i(x)}{\partial x^T} \right) = \exp \left( \sum_{j=1}^d s(x_1)_j \right)$$

# Real NVP

 **Meh...**

Let  $f = f_n \circ f_{n-1} \circ \dots \circ f_1$ , where  $f_i$  is defined as:

$$y_1 = x_1$$

$$y_2 = x_2 \odot \exp(s(x_1)) + t(x_1)$$

 **Meh...**

where  $(x_1, x_2)$  is some partition of the input.

- Easily invertible!
- Simple form of the determinant of the Jacobian:

$$\det \left( \frac{\partial f_i(x)}{\partial x^T} \right) = \exp \left( \sum_{j=1}^d s(x_1)_j \right)$$

# Instantaneous change of variables

Now, if we make the transformation continuous in time:

$$\frac{d\mathbf{z}}{dt} = f(\mathbf{z}(t), t)$$

we are also interested how log-probability changes in time. Instantaneous change of variable theorem shows that:

$$\frac{\partial \log p(\mathbf{z}(t))}{\partial t} = -\text{tr} \left( \frac{df}{d\mathbf{z}(t)} \right)$$

# Continuous Normalizing Flow (CNF)

1. **Training** (Data distribution  $\rightarrow$  Prior distribution):

$$\mathbf{z}_0 = \mathbf{z}_T + \int_T^0 f(\mathbf{z}(t), t) dt$$



# Continuous Normalizing Flow (CNF)

1. **Training** (Data distribution  $\rightarrow$  Prior distribution):

$$\mathbf{z}_0 = \mathbf{z}_T + \int_T^0 f(\mathbf{z}(t), t) dt$$

2. **Loss function:**

$$\log(p(\mathbf{z}_T)) = \log(p(\mathbf{z}_0)) - \int_0^T \text{tr} \left( \frac{df}{d\mathbf{z}(t)} dt \right)$$

# Continuous Normalizing Flow (CNF)

1. **Training** (Data distribution  $\rightarrow$  Prior distribution):

$$\mathbf{z}_0 = \mathbf{z}_T + \int_T^0 f(\mathbf{z}(t), t) dt$$



2. **Loss function**:

$$\log(p(\mathbf{z}_T)) = \log(p(\mathbf{z}_0)) - \int_0^T \text{tr} \left( \frac{df}{d\mathbf{z}(t)} dt \right)$$

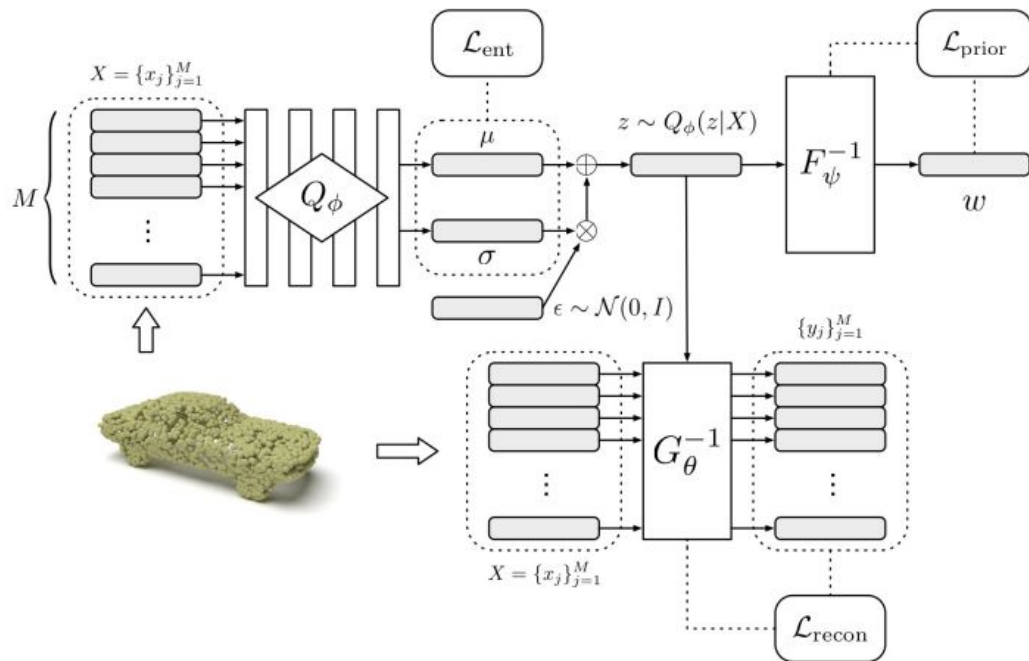
3. **Sampling** (Prior distribution  $\rightarrow$  Data distribution):

$$\mathbf{z}_T = \mathbf{z}_0 + \int_0^T f(\mathbf{z}(t), t) dt$$


# CNF example - FFJORD



# CNF example - PointFlow

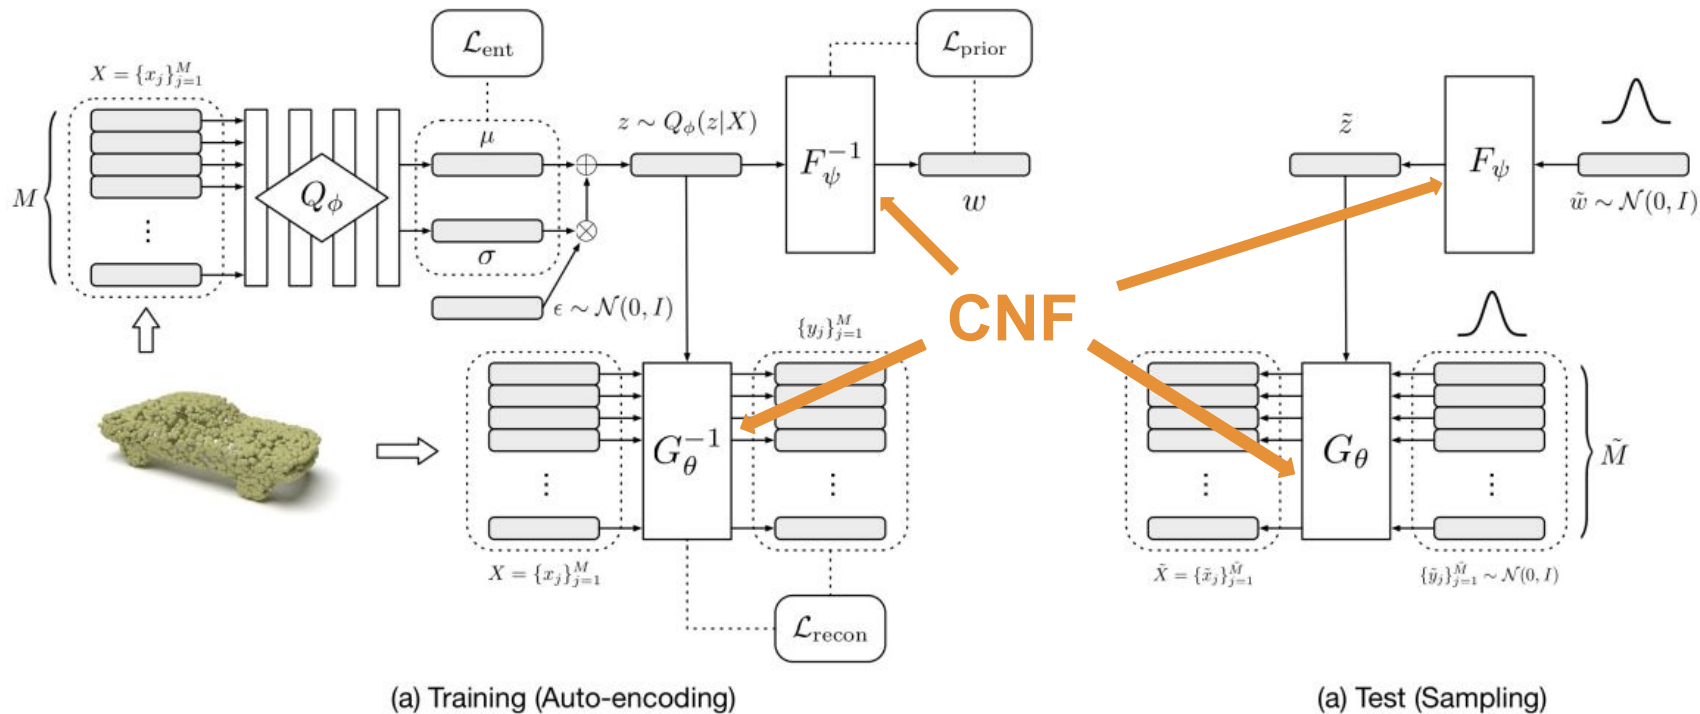


(a) Training (Auto-encoding)





(a) Test (Sampling)

# CNF example - PointFlow



# CNF example - PointFlow



Thank you!